

Contents

1	Executive summary	3
2	Introduction	4
3	Learning Bayesian networks	5
4	Software Design and Implementation	8
4.1	The core module	8
4.2	The learning module	10
5	Task 4.1: Parallelization of structural learning	13
6	Conclusion	16
	References	16

Document history

Version	Date	Author (Unit)	Description
v0.3	21/11 2014	Helge Langseth, Thomas D. Nielsen, Antonio Salmerón	First draft

1 Executive summary

The aim of this document is to describe the progress of the software development related to learning in the AMIDST project at the end of the first year.

[[Needs to be written]]

2 Introduction

[[This is just the first page from D3.1, then quickly terminated after learning is introduced. Just to say what a BN is.]]

Probabilistic graphical models provide a well-founded and principled approach for performing inference in complex domains endowed with uncertainty. A probabilistic graphical model is a framework consisting of two parts: a qualitative component in the form of a graphical model encoding conditional independence assertions about the domain being modelled as well as a quantitative component consisting of a collection of local probability distributions adhering to the independence properties specified in the graphical model. Collectively, the two components provide a compact representation of the joint probability distribution over the domain being modelled.

Bayesian networks (BNs) [?] are a particular type of probabilistic graphical model that has enjoyed widespread attention in the last two decades. Figure 1 shows a BN representing the joint distribution of variables X_1, \dots, X_5 . Attached to each node, there is a conditional probability distribution given its parents in the network, so that the joint distribution factorises as

$$p(X_1, \dots, X_5) = p(X_1)p(X_2|X_1)p(X_3|X_1)p(X_4|X_2, X_3)p(X_5|X_3).$$

In general, for a BN with n variables $\mathbf{X} = \{X_1, \dots, X_n\}$, the joint distribution factorises as

$$p(\mathbf{X}) = \prod_{i=1}^n p(X_i | \text{pa}(X_i)), \quad (1)$$

where $\text{pa}(X_i)$ denotes the set of parents of X_i in the network.




Figure 1: A Bayesian network with five variables.

We will use lowercase letters to refer to values or configurations of values, so that x denotes a value of X and \mathbf{x} is a configuration of the variables in \mathbf{X} . Given a set of observed variables $\mathbf{X}_E \subset \mathbf{X}$ and a set of variables of interest $\mathbf{X}_I \subset \mathbf{X} \setminus \mathbf{X}_E$, *probabilistic*

inference is the calculation of the posterior distribution $p(x_i|\mathbf{x}_E)$ for each $i \in I$. A thorough introduction to the state of the art for inference in Bayesian networks was given in [?].

In this document we assume that inference techniques for a given Bayesian network is available, and will consider how to define a Bayesian network model that fits a data set as well as possible. This process is known as *learning* in the Bayesian network community.

3 Learning Bayesian networks

Consider again the factorization of the full joint distribution $p(\mathbf{x})$ in Equation (1). With this factorization we can efficiently represent the joint probability distribution $p(\mathbf{x})$, but it requires that $\text{pa}(X_i)$, i.e., the *parent set* of X_i , is known for each $i = 1, \dots, n$. The parents of X_i are exactly the nodes which have an outgoing edge pointing to X_i (e.g., $\text{pa}(X_4) = \{X_2, X_3\}$ for the model in Figure 1). Algorithms for learning the parent sets, also known as *structural learning* algorithms, follow one of two approaches: *i*) search and score methods, see e.g. [?] and *ii*) constraint-based techniques [?]. Both will be considered in Task 4.1 of the project, and the progress towards the implementation of structural learning in AMIDST is summarized in Section 5.

As soon as the parent sets in Equation (1) are determined, the specific conditional distributions $p(X_i|\text{pa}(X_i))$ can also be learned from data. The approach taken in the AMIDST project is that the conditional distribution $p(X_i|\text{pa}(X_i))$ is assumed to be a member of the exponential family (including, among others, the Gaussian and multinomial distributions). Therefore, parameter learning amounts to finding the optimal *parameterization* of the given distributions. This task will be considered in Task 4.2 and Task 4.4.

Let θ denote the combination of all parameters of the model (that is, all parameterizations of the local conditional distribution functions in Equation (1)), and let $\mathcal{D} = \{\mathbf{x}_1, \dots, \mathbf{x}_N\}$ denote the dataset which we are learning from. \mathcal{D} is a collection of N independent configurations over the variables of the domain. Note that \mathcal{D} may be *incomplete*, in case some of the observations are missing. Missingness can occur, for instance, if some of the variables are latent, or some data-collecting instrument fails.

Now, parameter learning in Bayesian networks amounts to estimating θ using some optimization criterion that depend on \mathcal{D} . The learning generally comes in two different shapes. The most commonly used approach (at least traditionally) is *maximum likelihood* learning, which attempts to find the model parameters that maximizes the *likelihood* of the model given the data. The likelihood of the parameters given the data set \mathcal{D} , is defined as $\mathcal{L}(\theta|\mathcal{D}) = \prod_{j=1}^N p(\mathbf{x}_j|\theta)$, and the maximum likelihood estimator is $\hat{\theta} = \arg \max_{\theta} \mathcal{L}(\theta|\mathcal{D})$.

Maximum likelihood learning is very efficient in exponential family models when there are no missing observations in \mathcal{D} , because the parameters of each distribution can be

optimized independently. However, closed-form solutions are not available when the dataset contains missing values.

Missing values is, though, a key ingredient in the data sets entertained in the AMIDST project. In these cases, maximum likelihood learning typically relies on the Expectation Maximization (EM) algorithm or generalizations thereof. Implementations of the (generalized) EM algorithm iterates over the following two steps that are repeated until convergence: *i*) E-step: Inference in the model given the data-set (using a model with the current parameterization); *ii*) M-step: Updates of the parameters using the posterior distributions calculated in the E-step. The EM algorithm is a greedy algorithm, that ensures that the likelihood of the current parameter estimates is non decreasing from one iteration to the next, and therefore also converges to a (local) maximum of the likelihood function. Unfortunately, this is only guaranteed when an exact E-step is conducted, and if one uses approximative inference the algorithm may not improve as it moves along. Exact inference in the E-step is computational prohibitive in a streaming context; additionally, the M-step can also be computationally challenging. Maximum likelihood learning is therefore not implemented in the open-source software of the AMIDST project, but can be employed utilizing the interface that the toolbox offers to the implementation inside the Hugin system.

Alternatively to the maximum likelihood approach we have the *Bayesian* paradigm. From a simplistic point of view, the main difference between the two is that the Bayesian set-up regards θ as a vector of *random variables*, and treats them on an equal footing as all other (unobservable) variables in the domain. From a modelling perspective this extension is straight forward: The model in Figure 2 extends the model in Figure 1 by explicitly representing $\theta = \{\theta_i\}_{i=1}^5$, where θ_i is the parameterization of the conditional distribution function $p(X_i|\text{pa}(X_i))$. Parameter learning now amounts to calculating the probability distribution $p(\theta|\mathcal{D})$, that is, it is reduced to inference in an (extended) Bayesian network model. Thereby, approximate inference techniques like variational Bayes message passing, exponential propagation or importance sampling can be used directly for learning.

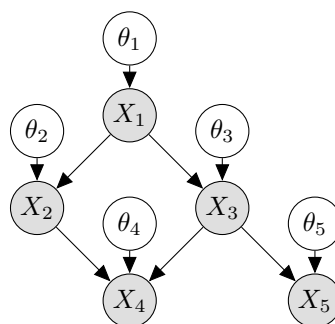


Figure 2: The Bayesian network from Figure 1 extended with explicit representation of the (unknown) parameters.

In practice, the *a priori* marginal distributions for each θ_i must be declared before the learning can be performed. These prior distributions enables domain experts to express knowledge about the parameterization of the distributions that is combined with information from the dataset to obtain the posterior information captured by $p(\boldsymbol{\theta}|\mathcal{D})$. The prior information consists of the definition of a distributional family for each θ_i together with a parameterization (the so-called *hyper-parameters*). In AMIDST we will ensure efficient calculation of posteriors by enforcing the distributional families be the conjugate distribution of the likelihood-terms. The hyper-parameters are chosen freely, and this is the vehicle provided to encode prior (expert) knowledge.

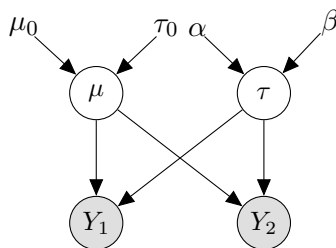


Figure 3: A detailed Bayesian network for the streaming variable Y_t (observed at time $t = 1$ and $t = 2$). Y_t is assumed to follow a Gaussian distribution with mean μ and precision τ .

Figure 3 shows a detailed description of the model for a streaming variable $Y_t \sim N(\mu, \tau^{-1})$ observed at $t = 1$ and $t = 2$. Prior information about these parameters are encoded by assuming $\mu \sim N(\mu_0, \tau_0^{-1})$ and $\tau \sim \Gamma(\alpha, \beta)$ for given values of the hyper-parameters $\{\mu_0, \tau_0, \alpha, \beta\}$. The dataset $\mathcal{D} = \{y_1, y_2\}$ enables us to calculate

$$p(\mu, \tau | y_1, y_2, \mu_0, \tau_0, \alpha, \beta) = \frac{p(y_1 | \mu, \tau) p(y_2 | \mu, \tau) p(\mu | \mu_0, \tau_0) p(\tau | \alpha, \beta)}{p(y_1, y_2 | \mu_0, \tau_0, \alpha, \beta)}.$$

The calculation is efficient both in terms of both space and time because the model is from the conjugate exponential family.

It follows that the learning functionality in AMIDST will rest heavily on the implementation of the core components in the toolbox (variables, distributions, Bayesian networks, and so on) as well the design to accommodate efficient inference using these core components. The implementation of inference engines will constitute a key component for the learning implementation because efficient and scalable *inference* is both a requirement and a guarantee for efficient and scalable *parameter learning*. The next section will therefore discuss the top-level design of the AMIDST toolbox, first describing the core components, thereafter moving to the design of the inference components.

4 Software Design and Implementation

In this section we provide an overview of the current design of the AMIDST toolbox as well as the status of the toolbox implementation. The first subsection gives a general overview of the core parts of the framework, which provide the foundation for the more learning specific aspects described in the second subsection.

4.1 The core module

An overview of the core components of the AMIDST toolbox is illustrated in Figure 4. The color coding in the figure summarizes the implementation status: blue boxes represent software components that have been implemented in the AMIDST toolbox and green boxes represent components that are part of the software design specification but which has not yet been implemented.

The structures included in the figure mainly relates to the basic components that play a key role in ensuring the forthcoming implementation of the different AMIDST learning and inference algorithms. For starters, instantiation of a particular probabilistic graphical model (PGM component) will be required. Currently it is possible to create either a static Bayesian network (BN component) or a two time-slice dynamic Bayesian network (2T-DBN component).

As previously described in Section 2, a BN consists of two components:

- A directed acyclic graph (DAG component) defined over a list of **Static Variables**. Each static variable is characterized by its name, ID, the state space type, the distribution type (i.e., multinomial or normal), as well as if it is observed or not.
- Conditional probability distributions of each variable given its parents (defined by the xDistributions component).

Similarly to a BN, a 2T-DBN (see Deliverable 2.1, Section 3.4) is defined using two main components:

- A dynamic acyclic graph (Dynamic DAG component) defined over a list of dynamic variables. The component specifies the graph structure of a 2T-DBN, i.e. the parent set for each variable at both time 0 and at time $t > 0$. Each dynamic variable is characterized by its name, ID, the state space type, the distribution type (i.e., multinomial or normal), and if it is observed or not. In order to represent the variables in a previous time step (needed when defining the dynamic DAG), we use the concept of *temporal clone* variables, which are copies of the real main variables but refer to the previous time step. For instance, X_{t-1} is codified as the *temporal clone* of variable X_t . Hence, in our data structures, the time index t is

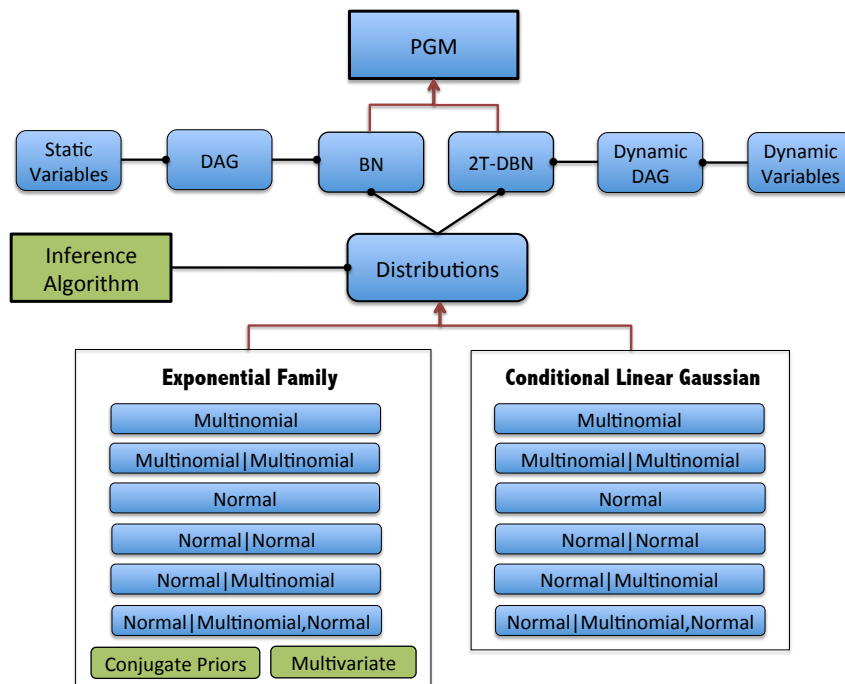


Figure 4: Illustration of the design of the software components related to core structures. Nomenclature: The boxes in the figure represent software components (sets, possibly singletons, of classes), a rounded-arc going from X to Y indicate that Y 'uses/references' X , and an arc with an arrow from X to Y implies inheritance.

not explicitly represented for a dynamic variable, but implicitly considered with the use of *temporal clones*.

- Conditional probability distributions of each variable given its parents (encoded in the Distributions component).

Note here that, in spite of the distinction between BN and 2T-BN, the distributions associated with the variables in both model classes can be defined in the same way. In particular, the Distributions component includes the set of conditional probability distributions considered in the AMIDST toolbox. More precisely, both variables with multinomial and normal distributions are modeled, and the distribution of each variable, in either a BN or a 2T-DBN, is initialized and specified according to its distribution type and the distribution types of its potential parents. This consequently gives rise to the following different implemented probability distributions:

- Multinomial: a multinomial variable with no parents
- Multinomial|Multinomial: a multinomial variable with multinomial parents.
- Normal: a normal variable with no parents.
- Normal|Normal: a normal variable with normal parents.
- Normal|Multinomial: a normal variable with multinomial parents.
- Normal|Multinomial,Normal: a normal variable with a mixture of multinomial and normal parents.

The case of a multinomial variable having normal parents is not considered yet in this initial prototype. It is planned to be included in future versions, although strongly restricted in inference and learning algorithms due to the methodological and computational issues previously commented in Deliverable 2.1.

We also provide an implementation of all the above distributions in the so-called Exponential Family form, which ensures an alternative representation of the standard distributions based on vectors of natural and moment parameters. This representation of the distributions support the type of parameter learning schemes pursued in AMIDST (see below).

4.2 The learning module

The core components of the modeling framework has been designed and implemented to facilitate an easy integration with the inference and learning modules developed as part of Work packages 3 and 4. Figure 5 shows a high-level overview of the key components of the AMIDST software tool that is directly related to learning; these learning-related

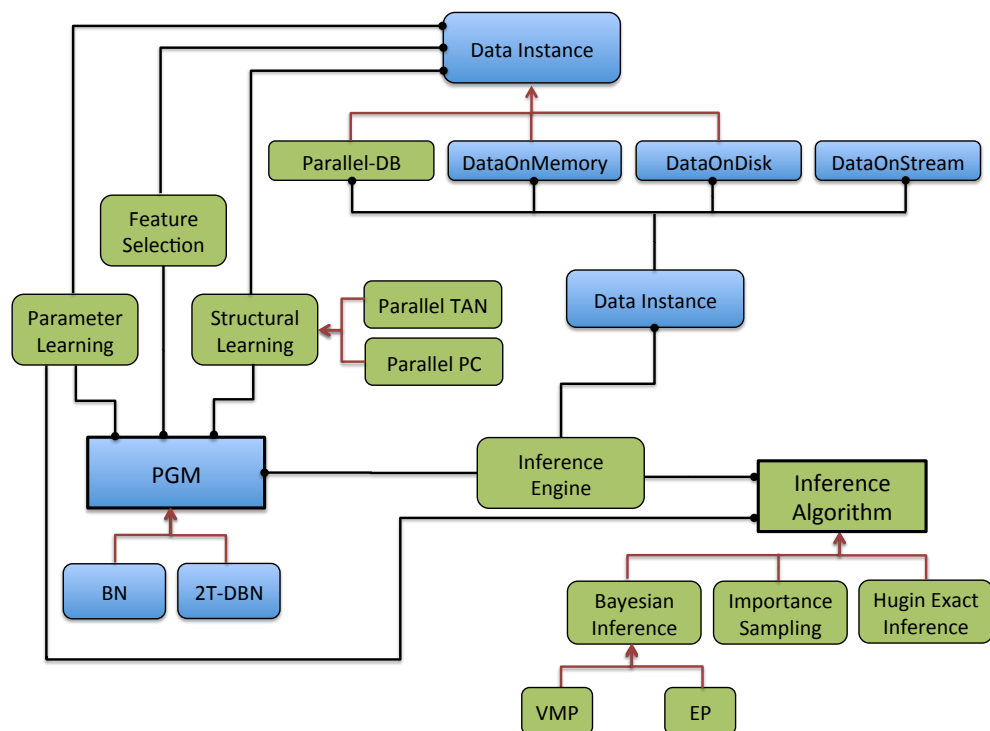


Figure 5: Illustration of the design of the software components related to model learning.

components are connected to the framework core (see Section 4.1) through the **Inference** component and the **PGM** component (shown using boxes with sharp corners).

In the AMIDST framework we consider two types of data sources for learning: i) Streaming data, where data arrives at high frequency with no storage of historical data (except for data in the most recent past, which is stored in a buffer), and 2) static databases that simply correspond to traditional databases. The database support is realized by a general database component (**Database**) that defines the database interface from which more specialized databases (**DataOnMemory**, **DataOnDisk**, and **ParallelDatabase**) can be derived:

- **DataOnMemory** implements database functionality for data sets that can be loaded into main memory.
- **DataOnDisk** provides functionality for handling datasets too large to be loaded in main memory.
- **ParallelDatabase** implements a distributed database.

The employed design is intended to support future users and developers of the AMIDST toolbox in the potential design and implementation of other database specifications; the only restriction being that new database components should implement the interface defined by the **Database** component.

Functionality for handling data streams is implemented by the **DataOnStream** component that allows data to be *pushed* to the AMIDST system (in contrast to a pull-approach that is standard when dealing with static database). To allow for variation in the run-time performance of the system a buffer component (**Buffer**) is needed for storing the most recent unprocessed cases. Each of the data sources are furthermore connected to the **Data Instance** component. This component consists of a single class that can represent a particular evidence configuration, such as the observed values of a collection of variables at time t or a particular row in a database.

Implementations of structural learning algorithms will be realized through the **Structural learning** component and its specialized sub-classes. The design includes components for supporting PC and TAN learning in a parallel setting (cf. Task 4.1). Currently, the corresponding implementation supports standard PC learning and parallel TAN learning by interfacing to the Hugin API.

As described in Section 3 we pursue a fully Bayesian approach for doing parameter learning in the AMIDST framework (cf. the activities in Task 4.2 and Task 4.4). This, in turn, means that parameter learning reduces to the task of inference for which we plan to consider two approaches: variational message passing (positioned in a variational Bayes framework and implemented in the **VMP** component) and expectation propagation (implemented in the **EP** component); both components are derived from the more general **Inference** component. Particular efficient implementations of both variational

inference and expectation propagation can be realized when the distribution families of the models are conjugate-exponential. In this case the inference operations can be further supported by specifying the exponential distributions using their natural parameters, see Section 3. These improvement are realized through tailored exponential family implementations of the standard distributions that are part of the AMIDST framework (such as the conditional linear Gaussian distribution). Note that the design of this parameter learning part of the overall AMIDST framework is flexible in the sense that it easily accommodates potential future learning-based extensions of the framework, e.g., Bayesian learning based on importance sampling or maximum likelihood learning using the expectation maximization algorithm (see also Section 3).

Variable selection is handled in the corresponding component in Figure 5, which is connected to both the model component (PGM) and the **Database** component.

The high-level design description above provides an overview of the current design of the AMIDST software framework with particular focus on the components that are directly related to model learning. The current status of the software implementation in relation to this design is as follows:

- The core components (marked as blue in Figure 4) has been implemented. This includes data structures for variables, graphs, Bayesian networks, dynamic Bayesian networks, key distributions such as multinomial and conditional linear Gaussian distributions represented in both standard form and as exponential families.
- Components defining data source management functionalities have been implemented, which includes support for handling static database (on disk and in memory) as well as streaming data.
- Methods for transforming AMIDST models to and from Hugin has been implemented based on the Hugin API.

Please see the appendix for the class diagrams for the implemented components.

Remark: The implementation is still ongoing and will continue so until the deliverable is submitted. Thus I expect that we will be able to expand the above list a bit. In particular, I know that there is currently work going on to provide parallel TAN learning in the AMIDST toolbox by calling the Hugin API. This would also be demonstrated at the review meeting.

5 Task 4.1: Parallelization of structural learning

Task 4.1 is devoted to the development of parallel algorithms for structural learning of Bayesian networks. The task covers *structure restricted models* as well as *classic constraint-based methods*.

Structure restricted models are sub-classes of Bayesian networks where only some particular structures are allowed. Typically, such kind of models are employed in specific tasks like *classification and regression* where one is interested in predicting the value of a target variable rather than in accurately modeling the dependencies among the variables in the model. A classification model contains a set of variables $\{X_1, \dots, X_n, C\}$ where C is the *class variable* and X_1, \dots, X_n are called *features*. A Bayesian network can be used for classification purposes by modeling the distribution $p(X_1, \dots, X_n, C)$. Then, an item with observed feature values x_1, \dots, x_n is classified as belonging to class c^* given by

$$c^* = \arg \max_{c \in \Omega_C} p(c|x_1, \dots, x_n) = \frac{p(c, x_1, \dots, x_n)}{p(x_1, \dots, x_n)}, \quad (2)$$

where Ω_C denotes the set of possible classes (i.e. the state space of variable C).

By restricting the possible structures for representing $p(X_1, \dots, X_n, C)$, it is possible to avoid the exponential growth in the number of parameters to learn from data with respect to the number of variables. The extreme case is when all the features are assumed to be independent given the class, which minimizes the number of parameters to learn, as the joint distribution factorizes as

$$p(X_1, \dots, X_n, C) = p(C) \prod_{i=1}^n P(X_i|C).$$

A classifier constructed in this way is called a naive Bayes (NB) classifier, and it corresponds to a Bayesian network structure as depicted in Fig. 6.

A n improvement on the NB is the so-called *Tree Augmented Naive Bayes* (TAN) classifier [?]. In a TAN, the features are arranged as a tree, and all the feature variables have the class as a parent. Figure 7 shows a TAN structure with five feature variables, X_1, \dots, X_5 .




Figure 6: A Naive Bayes structure

The structure of a TAN model is obtained by computing a *score* for each pair of features, namely their *conditional mutual information* given the class. Then, a maximum spanning tree of the features is constructed, labeling the edges using the computed scores.




Figure 7: A TAN structure. Unlike in Fig. 6, the features conform a tree rooted at X_2

Structural learning of TAN models has been approached in task 4.1 using two different methods. One of them uses *threads* to distribute the workload onto a number of cores. This method assumes that all data is available in main memory. The algorithm creates a number of threads that iterate through the scores to be computed. Each score is assigned to a thread and a thread only computes the scores assigned to it. The software developed under this approach corresponds to task 5.1 of WP5, and is included in the AMIDST toolbox inside HUGIN AMIDST (see Deliverable D5.1). It can be accessed from the open source AMIDST toolbox using the open source AMIDST \leftrightarrow HUGIN AMIDST interface, called HUGIN-Link, which is a functionality of the open source AMIDST toolbox that enables to use the HUGIN AMIDST API.

The other method uses *MPI* [?] to distribute the workload onto a number of processors. This method does not assume that all data is available in main memory. The variables are distributed between the processors and each process only reads the data for variables assigned to it. The computation of scores is controlled using Balanced-Incomplete Block Designs [?] as described in [?]. This approach assumes that data on each variable is stored in a separate file. This method is not available through the HUGIN-Link interface of the open source AMIDST toolbox. However, it can be accessed from it using file exchange, as the open source AMIDST toolbox is able to read and write HUGIN objects and files, and hence the structure yielded by HUGIN AMIDST can be transferred to software developed using the open source toolbox.

Constraint-based structural learning encompasses those algorithms for inducing unrestricted Bayesian network structures according to the result of a series of conditional independence tests. A brute-force constraint based algorithm could start off with a complete network (where each variable is linked to all the others) and remove edges between variables for which a statistical test accepts the independence hypothesis. Perhaps the most popular constraint-based structural learning algorithm is the so called PC [?]. Basically, the PC algorithm makes use of the same scores as the TAN, i.e. the conditional mutual information scores, which distribution is known to be of class χ^2 under the hypothesis of independence. It allows the construction of a statistical test for deciding about the independence relationships between the variables in the network at any given significance level.

Parallelization of the PC algorithm is still under development and the software will be implemented as part of task 5.1 of WP5. The functionality developed for implementing the parallel learning of TAN models will be used as a part of the parallel PC implementation. It will be available in the AMIDST toolbox as a part of HUGIN AMIDST and its access from the open source AMIDST will be effectively done in a similar way as described for the parallel TAN.

6 Conclusion

This document summarizes the progress of the software development related to learning functionality in the AMIDST toolbox. Parameter learning functionality rests heavily upon efficient design and implementation of core components and inference engines, and these parts were therefore discussed in some detail. Structural learning is ongoing (in Task 4.1) and will be finalized in Month 15. A preliminary discussion of the results was given. Implementation of dedicated functionality for parameter learning (Em algorithm and similar) has not yet started, and is not discussed.