

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
<b>2</b>	<b>Test and evaluation methodology</b>	<b>6</b>
2.1	Performance measures for classification on i.i.d. data . . . . .	6
2.1.1	Empirical risk . . . . .	8
2.1.2	Evaluation of families of classification rules . . . . .	9
2.1.3	Mann-Whitney $U$ test . . . . .	11
2.2	Performance measures for classification on streaming data . . . . .	11
<b>3</b>	<b>Cajamar: test and evaluation</b>	<b>12</b>
3.1	Use case requirements . . . . .	12
3.2	Model and data characteristics . . . . .	12
3.3	Predictive performance: test and evaluation . . . . .	13
3.3.1	Application scenario 1 . . . . .	13
3.3.2	Application scenario 2 . . . . .	13
3.4	Run-time performance: test and evaluation . . . . .	13
3.4.1	Application scenario 1 . . . . .	13
3.4.2	Application scenario 2 . . . . .	13
<b>4</b>	<b>Daimler: test and evaluation</b>	<b>13</b>
4.1	Use case requirements . . . . .	13
4.2	Model and data characteristics . . . . .	13
4.3	Predictive performance: test and evaluation . . . . .	14
4.3.1	Application scenario 1 . . . . .	14
4.3.2	Application scenario 2 . . . . .	14
4.4	Run-time performance: test and evaluation . . . . .	14
4.4.1	Application scenario 1 . . . . .	14
4.4.2	Application scenario 2 . . . . .	14
<b>5</b>	<b>Verdande: test and evaluation</b>	<b>14</b>

---

5.1	Use case requirements . . . . .	14
5.2	Model and data characteristics . . . . .	14
5.3	Predictive performance: test and evaluation . . . . .	15
5.3.1	Application scenario 1 . . . . .	15
5.3.2	Application scenario 2 . . . . .	15
5.4	Run-time performance: test and evaluation . . . . .	15
5.4.1	Application scenario 1 . . . . .	15
5.4.2	Application scenario 2 . . . . .	15
<b>6</b>	<b>Conclusion</b>	<b>15</b>

---

## Document history

Version	Date	Author (Unit)	Description
v0.3			The test and evaluation framework discussed and established
v0.6			Initial draft finished and reviewed by the PSRG
v1.0			Final version of document

# 1 Introduction

Even though the number of algorithms designed for learning on streaming data is increasing, there is still not a unified and well accepted way for evaluating them. This is because testing and evaluating algorithms that are designed to work on streaming data are generally more complicated than those designed to work on non streaming data. There are both statistical and computational reasons for this.

Static data is data, where each instance can be assumed to be identically and independently distributed i.i.d. If the data instances in a data stream are i.i.d., then the only new challenge compared to static data is that the data streams are open ended. It can be seen as a static data that continuously grow in size. One way to deal with this is to take out a data set of fixed size, a holdout data set, and perform tests on this. Another way is to let the performance measures be "accumulating" as the data set is growing.

Various performance measures on data streams, that have been proposed in the literature. Performance measures involving loss functions on stream data have been proposed in the papers of Gama et. al. [1], [2], [3]. The loss function on regression problems is a function of both the predicted value and the ground truth, while the loss function on classification problems is a function of predicted and real class labels. The holdout error is basically the average loss on a holdout dataset of fixed size. The predictive sequential, or *prequential* error is defined as the average loss function up to time step  $i$ , where  $i$  is the current time step.

Stationary data is data that follow a stationary process, i.e. that the joint probability distribution of any time window does not change when shifted in time. Consequently, parameters such as the mean and variance, if they are present, also do not change over time and do not follow any trends. It was suggested to use *prequential* error on stationary streams in [3]. Moreover, in the case when data was slowly drifting, it was suggested to use the prequential error measures with forgetting mechanisms. The forgetting mechanisms involves either using a time window or fading factors. In paper [3], convergence towards the Bayes error is shown for all these performance measures provided that the learners are consistent and data is i.i.d. Moreover, it is shown that in the case of concept-drift, i.e. when data is drifting and learners are changing over time, the prequential error measures with forgetting mechanisms are favorable.

In this paper, we focus the exposition on binary classification, although we are aware that much can be generalized to multi class classification or regression. Binary classification becomes complex when there is class *imbalance*, as known on non streaming data as well. Specialized learning algorithms on imbalanced streams are proposed in [4–6], where these papers points out that this problem is particularly difficult and effective algorithms for evaluating such a classifier is vital. In the papers [4, 6], the area under the receiver operator characteristics curve (AUC) is calculated on limited holdout sets, while in [5] AUC is calculated on the entire stream.

AUC is a popular method for evaluating classification problems where class imbalance is

---

vital, because it is invariant to the class distribution. AUC calculation requires that the classification rule is based on comparing a score or an output function to a threshold. In the Amidst software, this condition is often fulfilled as the output of the Bayesian network is commonly a predicted probability of the class.

AUC has the statistical interpretation that it is the probability that a member of the "positive" class is scored higher than a member of the "negative" class. AUC is therefore a measure of the *ranking ability* of the output function. This is particularly relevant if one wants to change the classification threshold as a consequence of changing class distributions or changing misclassification costs [7]. Moreover, it also pointed out that AUC is more preferable than accuracy for model evaluation in [8].

AUC calculation involves sorting all instances and iteration over the sorted list. In a streaming context, this means that the calculation is  $O(n)$  where  $n$  is the length of the data stream. If the whole stream is used, it may have computational problems related to memory and cpu time. Moreover, it is not taking into account concept-drift as discussed above. In [9], the prequential AUC with a forgetting mechanism is proposed. The advantages of this measure is that it allows for concept-drift and fixed memory and cpu time constraints.

In the AMIDST project there are use case scenarios where the data is multiple i.i.d. streams. These use case scenarios are dealt with by dividing these streams into sub streams over a particular time window, where a evaluation time is in between the start time and the end time. The data from the start time to the evaluation time is basically the explanatory variables and the data at the end time is basically the class label. This problem reduces to a problem of non streaming nature, and the evaluation is not taking concept-drift into account. However, it is worth noting that this problem can be expanded in the sense of prequential loss as of [3] or prequential AUC as of [9].

Moreover, there are also use case scenarios in the Amidst project that do not directly fit the prequential performance measures as of [3] and [9]. All of these measures involve calculating a loss or a score at each time step. In this report, we will describe a use case scenario where it is important to predict that certain event is happening somewhere inside an *interval* prior to the event actually takes place.

This problem has been solved by chopping the streams into a number of sub streams that are approximately i.i.d and labeled as either a positive or a negative. It is important that enough space between sub streams are left so that dependencies are withdrawn. In essence, the i.i.d. assumption of the sub streams allows one to investigate the output function independently from the prior distributions and the choice of loss functions.

In section 2, AMIDST relevant methodologies for evaluation of both batch and streaming algorithms are identified and discussed. This section forms the foundation of the next three sections, where the exact evaluation routines for each use case provider is given. These sections contains a description of the requirements related to evaluation as described in Delivery 1.2 and a short description of the algorithms and the data. At the end of these sections, the methods for evaluating predictive and runtime performances

---

		Predicted class		
		Cat	Dog	Rabbit
Actual class	Cat	5	3	0
	Dog	2	3	1
	Rabbit	0	2	11

Table 2.1: Example of a confusion matrix. A classifier is labelling instances as either cats, dogs or rabbits. The accuracy is the sum of the diagonal elements divided by the total number (in this case  $19/27$ ).

are exposed and discussed. Section 6 concludes the report.

## 2 Test and evaluation methodology

As discussed in the introduction, finding appropriate performance measures on streaming data is more difficult than finding performance measures on non streaming data. We will therefore start by discussing some relevant performance measures for classification on non streaming data that is i.i.d.

### 2.1 Performance measures for classification on i.i.d. data

We assume that the dataset has a fixed size  $n$ , where each instance is independently drawn from a joint probability distribution  $P(X, Y)$ , where  $X$  and  $Y$  are random variables. The  $X$  variable is known as the explanatory variable and  $Y$  is the class label. These two variables have output spaces  $\Omega_X$  and  $\Omega_Y$ , respectively. For instance in a binary classification problem  $\Omega_Y$  can either be true or false, while  $\Omega_X$  is a space of all possible explanatory vectors.

In classification, we typically consider a hypothesis function  $h : \Omega_X \rightarrow \Omega_Y$ . In terms of evaluating the performance of  $h$ , we use a dataset of  $n$  input-output pairs  $(x_i, y_i)$  that are independently drawn from  $X \times Y$ . The result of such an experiment can be shown in a confusion matrix. An example is shown table 2.1, where the classifier is attempting to distinguish cats, dogs and rabbits.

A global measure of the classification algorithm is the classification accuracy which is basically the diagonal elements divided by the total number (in this case  $19/27$ ). It is important to note that the accuracy is not telling the whole story of the classification rule. For instance, by looking at the table it is seen that the algorithm distinguishes cats from rabbits quite easily, but it is much harder to distinguish cats from dogs. This information can not be found by only looking at accuracy.

Moreover, accuracy is also very dependent on how evenly the classes are distributed.

		Predicted class	
		Cat	Not cat
Actual class	Cat	5	3
	Not cat	2	17

Table 2.2: Example of a confusion matrix for a classifier of cats and not cats. The accuracy is  $22/27$ .

		Actual Condition	
		Positive	Negative
Test outcome	Positive	5	2
	Negative	3	17

Table 2.3: Example of a confusion table for a classifier of cats and not cats. The true positives and true negatives are on the diagonal, while the two other numbers are the false positives and the false negatives.

For instance, when there are a lot more instances of one class compared to the others, a naive classification rule that always predict the majority class will get a high accuracy, even though the method is not using any of the information that is contained in the explanatory variables. These are reasons for inspecting the full confusion matrix for interpreting the classification rule. There are also more numbers that can be derived from the confusion matrix, but to simplify the exposition, we have chosen to limit the discussion to binary classification.

For a binary classifier, such as classifying cats and non-cats, the confusion matrix is two dimensional as shown in table 2.2. In binary classification, it is common to introduce positives and negatives, instead of the class labels. A true positive is therefore an actual cat that has been predicted to be a cat. False positives, true negative and false negatives are defined in an equivalent manner. The results are commonly shown in a confusion table (see table 2.3), which is not the same as a confusion matrix.

From a confusion table it is easy to calculate numerous numbers that describes the classification rule. Specifically, we mention the true positive and false positive rates. True positive rates, also known as recall, is the number of true positives divided by the total number actual positives. The false positive rates, also known as fall-out, is the number of false positives divided by the total number actual positives.

By investigating various numbers that can be deduced from the confusion table, it is possible to discuss classification rules, even when the datasets are unevenly distributed. That is, when some class label have more instances than others.

However, these numbers do not take into account that some misclassifications might be more costly than others. For instance, a false positive might be more costly than a false negative, such as in the case of cancer diagnostics. It might be more costly to

not treat a person that is sick, compared to treating a healthy person. Moreover, the cost of each false positive (or false negative) may not be constant either. For instance, if the classifier is predicting whether a client in a bank will default a loan or not, the cost is clearly related to the size of the loan in question. The next subsection includes a procedure to include such costs in a performance measure.

### 2.1.1 Empirical risk

In mathematical optimization, statistics, decision theory and machine learning, a loss function or cost function is a function that maps an event or values of one or more variables onto a real number that is intuitively representing some *cost* associated with the event. Loss functions can be used on optimization problems, where an algorithm or method is optimized by minimizing the loss function. Moreover, loss functions are frequently used to diagnose and compare various algorithms or methods.

In this paper, we define the *loss function* as a real and lower-bounded function  $L$  on  $\Omega_X \times \Omega_Y \times \Omega_Y$ . The value of the loss function at an arbitrary point  $(x, h(x), y)$  is interpreted as the loss, or cost, of taking the decision  $h(x)$  at  $x$ , when the right decision is  $y$ . Notice that in this paper, the loss function is dependent on  $x$  as well. This is of high practical use, because a certain misclassification might be more expensive than another.

In the frequentist perspective, the expected loss is often referred to as the risk function. It is obtained by taking the expected value over the loss function with respect to the probability distribution  $P(X, Y) : \Omega_X \times \Omega_Y \rightarrow \mathbb{R}^+$ . The *risk function* is given by

$$R(h) = \int_{\Omega_X, \Omega_Y} L(x, h(x), y) dP(x, y). \quad (2.1)$$

In the case when the costs are independent of  $x$  and also that there is no cost related to correct classification, the risk function reduces to the well known expected cost of misclassification (ECM)

$$ECM = c(1|0)p(1|0)p_0 + c(0|1)p(0|1)p_1. \quad (2.2)$$

Here,  $c(1|0)$  is the cost for misclassifying an item of class zero as class one and  $p(1|0)$  is the misclassification probability given class zero. The quantities  $c(0|1)$  and  $p(0|1)$  are defined equivalently, while  $p_0$  and  $p_1$  are the priors.

In general, the risk  $R(h)$  cannot be computed because the distribution  $P(x, y)$  is unknown. However, we can compute an approximation, called empirical risk, by averaging the loss function on the training set  $\mathbf{x} \times \mathbf{y}$  of size  $n$ , where each element is  $(x_i, y_i)$ . The empirical risk is given by



$$R_{emp}(h, \mathbf{x} \times \mathbf{y}) = n^{-1} \sum_{i=1}^n L(x_i, h(x_i), y_i). \quad (2.3)$$

Notice that  $L$  is an array of  $n \times 2 \times 2$  elements. Many supervised learning algorithms are optimized by finding the  $h$  in a hypothesis space  $\mathcal{H}$  that minimizes the empirical risk. This paper will not focus on empirical risk minimization, but rather focus on using the empirical risk to compare methods.

### 2.1.2 Evaluation of families of classification rules

So far we have discussed how to evaluate a single classification rule. However, most classification rules in the AMIDST framework is based on comparing an estimated probability to a certain threshold. We call this estimated probability the output function  $q : \Omega_X \rightarrow \mathbb{R}^+$ . In the AMIDST framework the range of  $q$  is usually  $[0, 1]$ , but this restriction is not necessary for this theory. The potential classification rules are the family of hypothesis functions  $\mathcal{H}$ , where each element  $h_T : \Omega_X \rightarrow \Omega_Y$  has the form

$$h_T(x) = \begin{cases} 0 & \text{for } q(x) \leq T \\ 1 & \text{else.} \end{cases} \quad (2.4)$$

It is of interest to evaluate all these classification rules. The receiver operating characteristic ROC is a plot of the true positive rate as a function of the false positive rate, as  $T$  vary over all relevant variables. ROC analysis provides tools to select possibly optimal models and to discard suboptimal ones independently from (and prior to specifying) the cost context or the class distribution. ROC analysis is related in a direct and natural way to cost/benefit analysis of diagnostic decision making.

The ROC curve allows a visual exposition of the family of classifiers  $\mathcal{H}$ . In particular, it is easy to see which threshold is needed for a certain hit rate or true positive rate. Moreover, it is also of interest to reduce all this information into a single number. The area under the ROC curve, also known as AUROC, is such a number. AUROC is independent of  $T$  and also the priors. It may be difficult to comprehend what AUROC is at this point, but we will return to it afterward.

First, let  $X_0$  and  $X_1$  be random variables with probability distributions  $P(X|Y = 0)$  and  $P(X|Y = 1)$ , respectively. We define the random variables  $Q_0 = q(X_0)$  and  $Q_1 = q(X_1)$ . To summarize,  $Q_0$  is basically the output value you get when you pick a random sample of class zero. The same thing can be said for  $Q_1$ .

The probability  $P(Q_1 > Q_0)$  is of interest, because this says what is the probability that if you take one sample from each of the populations, what is the chance that the output value from sample one is higher than the output value of sample zero. This question is independent of  $T$ , meaning that the discussion about priors and costs are not needed. This probability is called the concordance probability.




Figure 2.1: Receiver operating characteristics curve. True positive or hit rate is shown as a function of false positive or false alarms rate. The dashed line shows the ROC curve of the random guess.

Without exposing the details, it is possible to show that

$$\text{AUROC} = P(Q_1 > Q_0). \quad (2.5)$$

It is important to note that nothing need to be assumed about the probability distributions of  $Q_0$  and  $Q_1$ . Furthermore, it is worth noting that the concordance probability is exactly equal the common language effect size of the Mann-Whitney  $U$  test. We have therefore chosen to add a section of the Mann-Whitney  $U$  test.

### 2.1.3 Mann-Whitney $U$ test

In statistics, the Mann-Whitney  $U$  test (also called the Mann-Whitney-Wilcoxon (MWW), Wilcoxon rank-sum test, or Wilcoxon-Mann-Whitney test) is a nonparametric test of the null hypothesis that two populations are the same against an alternative hypothesis, especially that a particular population tends to have larger values than the other. It has greater efficiency than the  $t$ -test on non-normal distributions and it is nearly as efficient as the  $t$ -test on normal distributions.

We define a training set  $\mathbf{x} \times \mathbf{y}$  with  $n$  input-output pairs  $(x_i, y_i)$ , independently drawn from  $P(X, Y)$ . From the training set we have two populations  $\mathbf{q}_0 = \{q(x_i), | y_i = 0\}$  and  $\mathbf{q}_1 = \{q(x_i), | y_i = 1\}$ . Their sizes are  $n_0$  and  $n_1$  so that  $n_0 + n_1 = n$ . Calculating the  $U$  statistics is straightforward, where these two values are obtained

$$U_0 = \sum_{i=1}^{n_0} \sum_{j=1}^{n_1} H(q_j - q_i) \quad \text{and} \quad U_1 = \sum_{i=1}^{n_0} \sum_{j=1}^{n_1} H(q_i - q_j). \quad (2.6)$$

Here,  $H(\cdot)$  is the heaviside step function and notice that  $U_0 + U_1 = n_0 n_1$ . For large samples, each  $U$  is approximately normally distributed. In that case, the standardized value

$$z = \frac{U_0 - m_U}{\sigma_U}, \quad (2.7)$$

where  $m_U$  and  $\sigma_U$  are the mean and standard deviation of  $U$  given by

$$m_U = \frac{n_0 n_1}{2} \quad \text{and} \quad \sigma_U = \sqrt{\frac{n_0 n_1 (n_0 + n_1 + 1)}{12}}. \quad (2.8)$$

Significance of test can be checked in tables of the normal distribution. Although, such an hypothesis test is interesting by itself, we are more interested in the concordance probability  $P(Q_1 > Q_0)$  which is defined by

$$P(Q_1 > Q_0) = \frac{U_1}{n_0 n_1}. \quad (2.9)$$

It is important to note that even though the Mann-Whitney  $U$  test is dependent on whether the shapes of the probability distributions of  $Q_0$  and  $Q_1$  are similar, this requirement is not necessary for the concordance probability. Also, equation (2.9) shows a simple way of calculating  $P(Q_1 > Q_0)$  and AUROC.

## 2.2 Performance measures for classification on streaming data

In the use case scenarios in the Amidst software, all problems are binary classifiers, which involve dividing the streams into sub streams that are labeled as either positives or negatives. There are two important approaches here.

The first approach deals with multiple streams and the global time dependence is removed by sampling all these streams into multiple sub streams, which are starting and ending at the same time steps. Each sub stream is labeled as either a positive or negative based on the data at the end time. An evaluation time step that is a fixed between the start and the end is also defined. The explanatory variables are derived from the data between the start and the evaluation time step. In this sense, the data is basically i.i.d.

The second approach is when there is basically one stream where one can assume stationarity on a global scale. Sub streams with large enough distance to other sub streams are taken out of the stream. These sub streams are labeled as either a positive or a negative. For instance, a positive is when something is happening towards the end of

the sub stream, while a negative is a sub stream where nothing happens. When the AMIDST software run, the output function is evaluated at each time step. Depending of the problem, a heuristic algorithm is chosen to output a single number for each sub stream based on the output values. This can for instance be the maximum value of the output function on an interval that is prior to an event.

### 3 Cajamar: test and evaluation

#### 3.1 Use case requirements

*Summarize the use case requirements for the different application scenarios. This information should be derived from Deliverable 1.2.*

#### 3.2 Model and data characteristics

*Describe aspects of the model and data relevant for the ensuing test and evaluation discussion. Much of this information can be synthesized from the existing documents, and should serve to make the document more self-contained.*

#### 3.3 Predictive performance: test and evaluation

##### 3.3.1 Application scenario 1

##### 3.3.2 Application scenario 2

#### 3.4 Run-time performance: test and evaluation

##### 3.4.1 Application scenario 1

##### 3.4.2 Application scenario 2

### 4 Daimler: test and evaluation

#### 4.1 Use case requirements

*Summarize the use case requirements for the different application scenarios. This information should be derived from Deliverable 1.2.*

---

## 4.2 Model and data characteristics

*Describe aspects of the model and data relevant for the ensuing test and evaluation discussion. Much of this information can be synthesized from the existing documents, and should serve to make the document more self-contained.*

## 4.3 Predictive performance: test and evaluation

### 4.3.1 Application scenario 1

### 4.3.2 Application scenario 2

## 4.4 Run-time performance: test and evaluation

### 4.4.1 Application scenario 1

### 4.4.2 Application scenario 2

## 5 Verdande: test and evaluation

### 5.1 Use case requirements

*Summarize the use case requirements for the different application scenarios. This information should be derived from Deliverable 1.2.*

### 5.2 Model and data characteristics

*Describe aspects of the model and data relevant for the ensuing test and evaluation discussion. Much of this information can be synthesized from the existing documents, and should serve to make the document more self-contained.*

---

### 5.3 Predictive performance: test and evaluation

#### 5.3.1 Application scenario 1

#### 5.3.2 Application scenario 2

### 5.4 Run-time performance: test and evaluation

#### 5.4.1 Application scenario 1

#### 5.4.2 Application scenario 2

## 6 Conclusion

## References

- [1] Gama, J.a., Rodrigues, P.P., Sebastião, R.: Evaluating algorithms that learn from data streams. In: Proceedings of the 2009 ACM Symposium on Applied Computing. SAC '09, New York, NY, USA, ACM (2009) 1496–1500
  - [2] Gama, J., Sebastião, R., Rodrigues, P.P.: Issues in evaluation of stream learning algorithms. In: 15th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD). (2009) 329–337
  - [3] Gama, J., Sebastião, R., Rodrigues, P.P.: On evaluating stream learning algorithms. Machine Learning **90**(3) (2013) 317–346
  - [4] Ditzler, G., Polikar, R.: Incremental learning of concept drift from streaming imbalanced data. IEEE Transactions of Knowledge and Data Engineering **25**(10) (2013) 2283–2301
  - [5] Hoens, T., Chawla, N.: Learning in non-stationary environments with class imbalance. In: 18th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD). (2012) 168–176
  - [6] Lichtenwalter, R., Chawla, N.: Adaptive methods for classification in arbitrarily imbalanced and drifting data streams. In: New Frontiers in Applied Data Mining. Volume 5669 of Lecture Notes in Computer Science. (2010) 53–75
  - [7] Wu, S., Flach, P., Ramirez, C.: An improved model selection heuristic for auc. In: ECML PKDD. Volume 4701 of Lecture Notes in Computer Science. (2007) 478–489
  - [8] Gama, J.: Knowledge Discovery from Data Streams. Chapman and Hall (2010)
-

- [9] Brzezinski, D., Stefanowski, J.: Prequential auc for classifier evaluation and drift detection in evolving data streams. In: Proceedings of the 3rd International Workshop on New Frontiers in Mining Complex Patterns, Nancy, France, September 19, 2014. (2014)
-